

PG4200: Algorithms And Data Structures

Lesson 06: Hash Maps and Sets

Bogdan Marculescu

Hash Function


- A function that maps data from an arbitrary size to a specific size
 - eg, mapping strings to a int
- $h(x)=y$, mapping from domain X to a value in domain Y
- $|X|$ is often much larger than $|Y|$
 - Note: in mathematics, given a set X , with $|X|$ we represent its cardinality, ie the number of elements in it

Hash Properties

- *Deterministic*: for a given input x' , should always get the same output y'
- *Uniform*: mapping from X to Y should be ideally spread uniformly over Y ,
 - ie the number of elements in X that map to a specific y' should be close to $|X|/|Y|$
- *Performance*: either fast (in this course) or slow (security, eg hashing of passwords)

Collisions

- If $|X| > |Y|$, you cannot avoid $h(x')=h(x'')$, two different values in X mapping to the same value in Y
- Ideally, if uniform, no more than $|X|/|Y|$ collisions per element



Hash Function Examples

- int to int
 - same domain, so same size
- long to int
 - from 2^{64} to 2^{32} values, which means that for every single int, there are $2^{64}/2^{32} = 2^{32}$, ie 4 billion, longs
- String to int
 - considering that a Java String can be composed of up to 2 billion characters (max length of its internal **byte[]** array), the space of all possible strings is astronomically huge...

```
public static int hash(int x) {
```

```
    return 1;
```

```
}
```

```
public static int hash(int x) {  
  
    return 1;  
}
```

- Not particularly good...
- It is NOT *uniform*, as all values are mapped to the same value 1

```
public static int hash(int x) {  
    return x;  
}
```

- So called *identity* function
- Technically, it is a valid hash function (eg, deterministic and uniform)
- Potential issue: trivial to revert, ie, knowing output, it is trivial to find the input
 - this would had be a huge problem if we were in the context of **security**, eg hashing of passwords, but not here in this course


```
public static int hash(int x) {  
    return x + 100;  
}
```

- Still a valid hashing function
 - the addition overflow for high values is irrelevant here
- But still trivial to revert, eg find input given output
- As in this course we do not deal with security, we are mainly interested in hash functions from larger to smaller sets

```
public static int hash(long x) {  
    return (int) x;  
}
```

- What does it mean to “*cast*” a long to a int?
- long is 64 bits, whereas int is just 32
- Somehow, here we lose information



- Ignore first 32 bits in the **long**, and take last 32 to make an **int** out of them
- WARNING: if the first bit in the right part is a 1, then the resulting **int** will be negative
 - as first bit in a **int** is used for the sign




- Issue: any change in a **long** in its left part will end up in exactly the same hash value
- What if we want to guarantee that a single bit change must result in a different hash?

```
public static int hash(long x) {  
    return (int) (x ^ (x >>> 32));  
}
```

- If, looking at it, your first reaction is “WTF?!?” do not worry... it is a normal reaction...
- \wedge : xor operator
 - WARNING: usually in mathematics the symbol \wedge represents an exponent, eg, 2^32 in mathematical notation is 2^{32}
- \ggg : right-shift operator

$X \ggg 32$



- Move all bits to the right by 32 positions
- On the left, empty positions will be filled with 0s
- On the right, old elements are discarded
- Note: there is a difference between “>>>” and “>>”
 - “>>>” ignores the sign, and always fill left part with 0s

`(int) (x ^ (x >>> 32))`

- Xor \wedge :
 - $0 \wedge 0 = 0$
 - $0 \wedge 1 = 1$
 - $1 \wedge 0 = 1$
 - $1 \wedge 1 = 0$



XOR \wedge



- Changing any single bit in either the left or right part will guarantee that the resulting xor will be different
- Note: changing more than 1 bit does not guarantee it... we are still going to have 4 billion collisions per `int`
- Recall that with `(int)` we discard the left part of the resulting long value

```
public static int hash(String x) {  
    int sum = 0;  
    for (int i = 0; i < x.length(); i++) {  
        sum += x.charAt(i);  
    }  
    return sum;  
}
```

- A **String** could have any size...
 - actually bounded by its internal **byte[]** representation, where sizes of arrays are **int**, so at most 2GB of data
- We can look at each **char** as 16-bit number, and sum all of them
 - don't care of integer overflow


```
public static int hash(String x) {  
    int sum = 0;  
    for (int i = 0; i < x.length(); i++) {  
        sum += x.charAt(i);  
    }  
    return sum;  
}
```

- Problem: many small strings will end up with the same hash, especially permutations
- eg, `hash("ab") = hash("ba")`
- eg, `hash("B") = hash("!!")`
 - B has numeric code 66, whereas ! has code 33

```
public static int hash (String x) {  
    final int delta = 31;  
  
    int sum = 0;  
    for (int i = 0; i < x.length(); i++) {  
        sum = (delta * sum) + x.charAt(i);  
    }  
    return sum;  
}
```

- Multiply by a constant at each character
- We are still going to have collisions, but less between short strings using Latin letters (which are the ones we usually work with)

Hash Maps

- Still a map from a K key to a V value
- No requirement on ordering of K keys, just being able to compute an *hash* of it
- In Java, all objects inherits from *java.lang.Object*, which defines a *hashCode()* method
- Hash code used as an index for an internal array

Example

- $put("foo", v)$
- $h("foo")=42$
- $h("foo") \% 10 = 2$
- Benefit: operations (insert/search/etc) have cost due to hash independent of size N of the collection


Internal array buffer of size $M=10$

[illegible]

What About Collisions?

- $put("foo", v)$
- $put("bar", z)$
- $h("foo") = h("bar")$
 - ie, collision due to same hash
- $h("foo") \% 10 = 2$
- What to do?

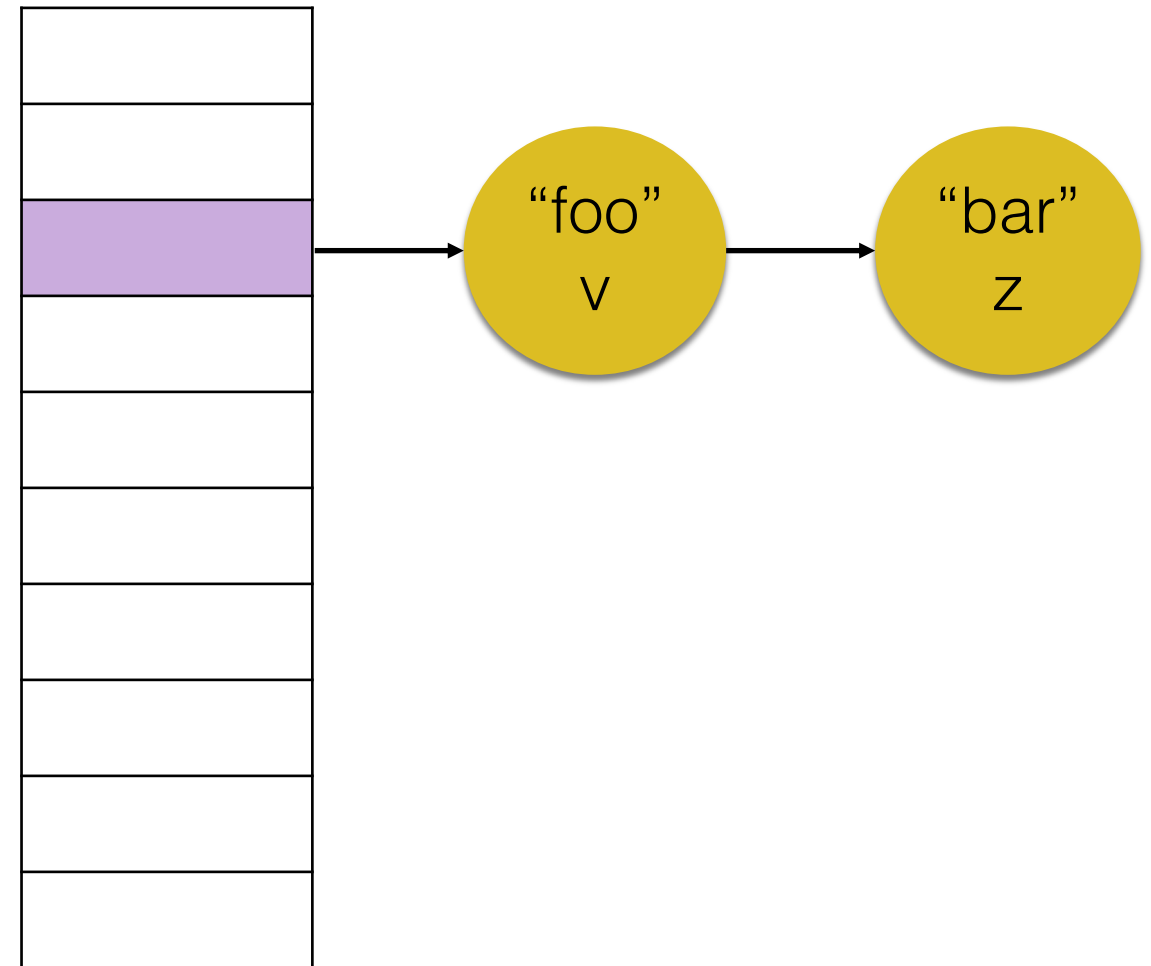
Internal array buffer of size $M=10$




Different Strategies

- $put("foo", v)$
- $put("bar", z)$
- $h("foo") = h("bar")$
 - ie, collision due to same hash
- Use list at each position sharing same hash
- Nodes containing keys and values

Internal array buffer of size $M=10$



hashCode() computed on the keys to determine their bucket.
In this example, assuming
“foo”.hashCode() == “bar”.hashCode(), because same bucket.
However, *“foo”.equals(“bar”)* is false



Once the bucket is determined with *hashCode()*, we use *equals()* on the *keys* in the list (one at a time), to see if there is a match

java.lang.Object

- *Object* does define two methods: *hashCode()* and *equals()*
- Those methods will depend on the internal fields of the object
- *Important*: if two objects are equals, then they **MUST** have same hash code
 - $A.equals(B)$ implies $A.hashCode() == B.hashCode()$
 - The vice-versa is not necessarily true, ie $A.hashCode() == B.hashCode()$ does not imply $A.equals(B)$, although that could happen
- What if constraint is not satisfied? Expect weird bugs when using maps and sets...

Cost

- Worst case: $O(N)$ if all elements end up in same “bucket” (ie same value for $h() \% M$), the map would be equivalent to a list
 - operations to search on list would be $O(N)$, albeit insert would be $O(1)$
- But, if M large enough compared to N , and hash function is uniform enough, you can have a $O(1)$ cost in many cases
 - even if you have some collisions, it will not be a problem, as you would have a small number of elements in the list

Hash or RBT?

- Hash Maps is the most popular and widely used
- If you know how much data you'll insert at most, can choose a good large enough M
- So in most cases, we are in $O(1)$ Hash vs $O(\log N)$ RBT
- But Hash can be $O(N)$ in worst case, vs RBT guarantees $O(\log N)$ in all cases
 - eg, in critical systems where you MUST guarantee a response within a certain amount of time, might want to use RBT
- Hash does not need ordering of keys

Set

- In mathematics, a *set* is a collection of elements where:
 - 1) *ordering is not important*: ie $\{1,2,3\}$ is equivalent to $\{2,3,1\}$
 - 2) *no repetitions*: ie $\{1,2\}$ is the same as $\{2,1,1,2,2,1,1,2,1\}$
- How to implement a Set in Java?
- Easy: use an internal $Map<K, V>$ where your values in the set are the keys K , and you just ignore the values V

Keys and Immutability

- *Immutable Object*: an object whose state cannot be changed once created
- Example: Strings are immutable
 - eg, concatenation with + and methods like *toUpperCase()* and *substring()* do NOT change the String, but rather *create* a NEW one
- Keys in a Map/Set **MUST** be *immutable*... why?

Different Hash

Foo foo = new Foo();

set.add(foo);

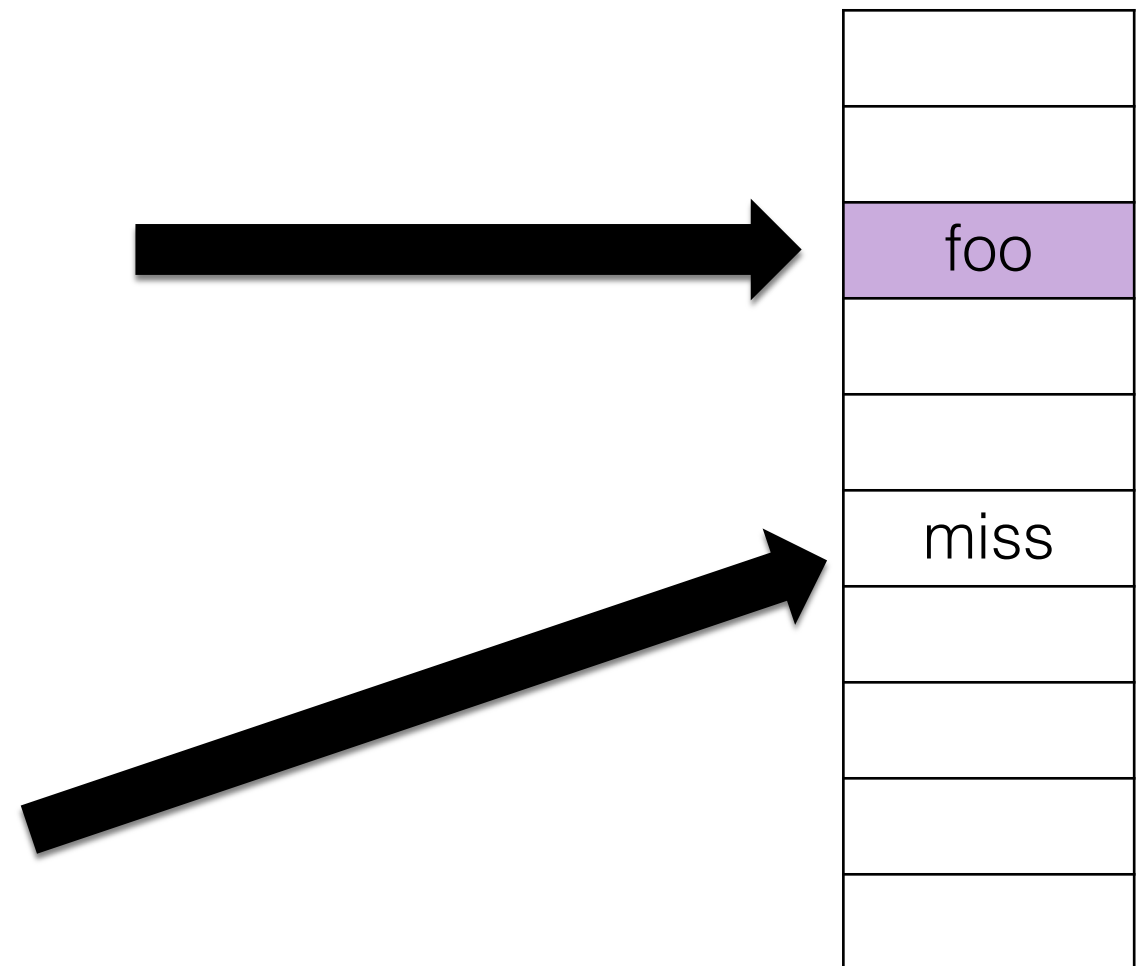
assertTrue(set.contains(foo));

// $h(foo) = 42$, $42 \% M = 2$

foo.setSomeVariable(...);

// $h(foo) = 55$, $55 \% M = 5$

assertFalse(set.contains(foo));



Using Maps and Sets

- Should only use a *Set* for immutable types
- What if you need a collection of mutable types $\langle X \rangle$?
 - in most cases, creating a *Set* $\langle X \rangle$ would be wrong!
- Option 1: rather use a list, eg *List* $\langle X \rangle$
 - however, it would allow duplicates
- Option 2: use a map *Map* $\langle K, X \rangle$ where the key is an immutable field derived from X
 - eg, if mutable *User*, *map.put*(*user.getId()*, *user*), where the id could be a String (recall strings are immutable)

Creating an Immutable Object

```
public class UserImmutable {  
  
    private final String name;  
    private final String surname;  
    private final int id;  
}
```

- The first step is to make each field “*final*”
 - this means that they cannot be modified
- This is enough for primitive values (eg, int and double) but not for objects!
 - those must be immutable as well!!!

“final” Limitation

```
final int[] foo = {1,2,3};
```

```
foo = {4,5}; // would not compile
```

```
foo[0] = 42; // no problem!
```

- A “final” reference cannot be modified, but you can modify the referenced state!
- Note: no problem for “*final String*”, as String is immutable
- So, if you have a mutable object **M** inside an immutable one **X**, must guarantee that no operation in **X** can change the state of **M** once initialized

Modifying an Immutable Object

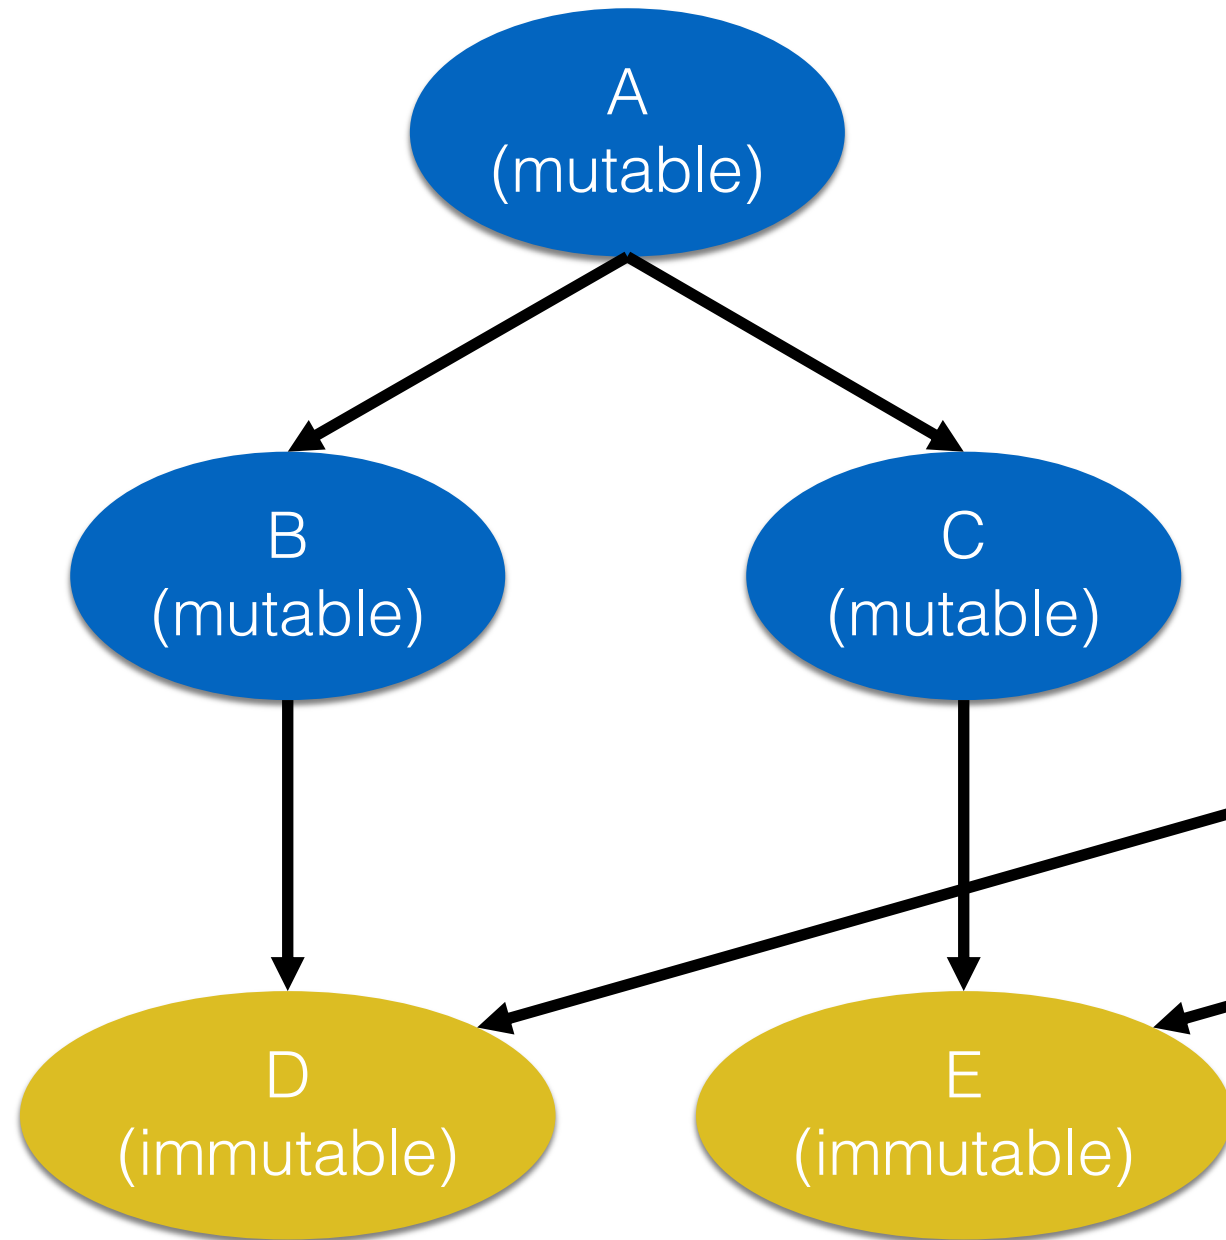
```
public UserImmutable withName(String s){  
    return new UserImmutable(s, surname, id);  
}
```

- Technically, you cannot modify an immutable object
- However, can return a new instance with the modified state
- Recall all methods in String return a new copy of the String, as String is immutable

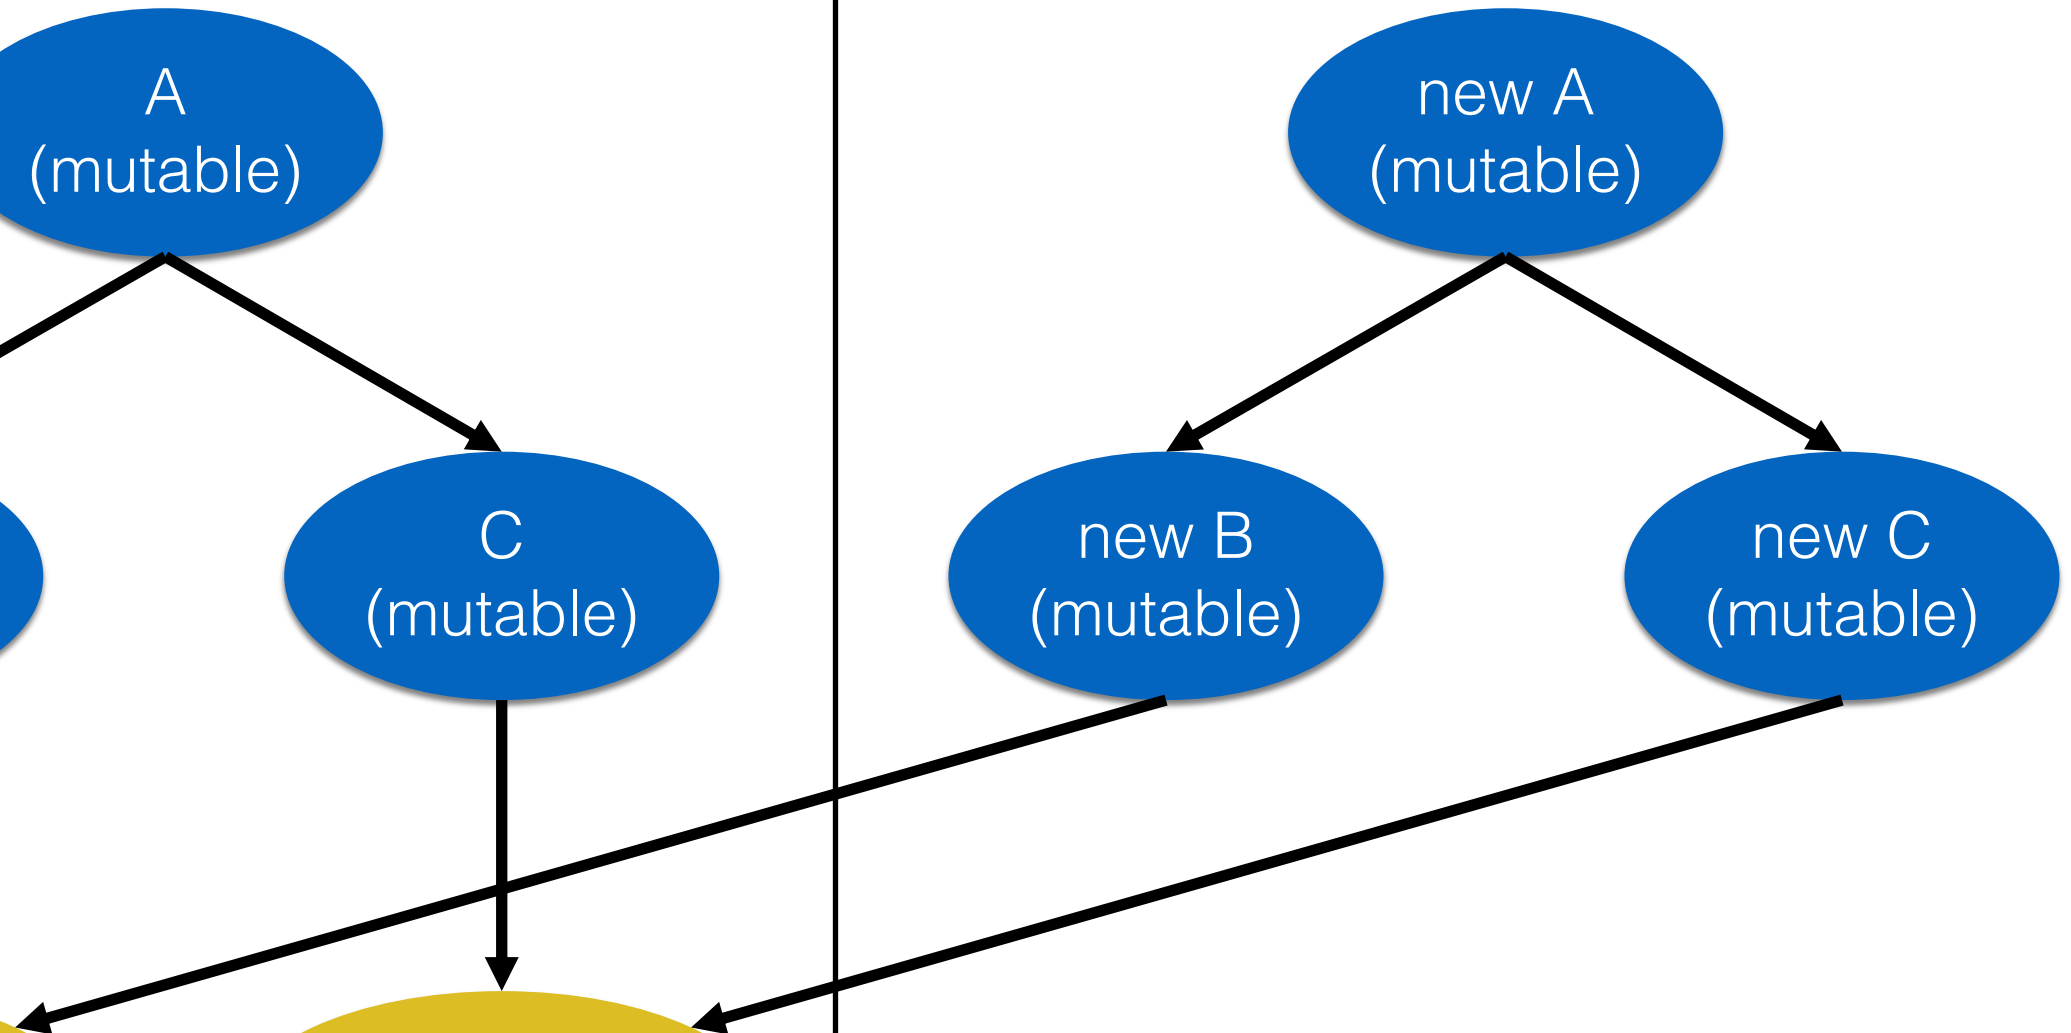
Object Copy

- When copying an object, it is a huge difference on whether its fields are **mutable** or **immutable** objects
- **Immutable**: can just copy the reference, ie, actual objects can be shared
 - aka “*shallow copy*”
- **Mutable**: must make a **new** complete copy, with **new** memory allocation on the heap
 - aka “*deep copy*”

Original



Copy



Homework

- Study Book Chapter 3.4 and 3.5
- Study code in the *org.pg4200.les06* package
- Do exercises in *exercises/ex06*
- Extra: do exercises in the book